

The Central Limit Theorem: The Mathematical Pillar of Gaussian Convergence and Statistical Inference

NECKER

June 5, 2026

Contents

1	Introduction	2
2	Formal Mathematical Statement	2
3	The Deep Meaning of the "Gaussian Collapse"	2
3.1	Why does the variance shrink? (σ^2/n)	3
4	Graphical Visualization of Convergence	4
5	Practical Example: Rolling Dice	4
6	What Is It Used For in Practice?	5

1 Introduction

In the real world, natural, social, or industrial phenomena follow very different distributions: waiting times at a counter follow an exponential distribution (highly asymmetrical), rolling a die follows a uniform distribution (flat), and population incomes often follow log-normal distributions.

The **Central Limit Theorem (CLT)** establishes that, regardless of the chaotic, asymmetrical, or bizarre shape of the starting population distribution, if we begin to extract samples and compute their mean, **the distribution of these sample means will inexorably tend to collapse toward a perfect geometric shape: the Gaussian curve (or Normal distribution).**

2 Formal Mathematical Statement

Let $X_1, X_2, \dots, X_n$ be a sequence of n independent and identically distributed (*i.i.d.*) random variables, drawn from a population with:

- Known and finite Mean (Expected Value): $E[X_i] = \mu$
- Known and finite Variance: $Var(X_i) = \sigma^2$

We define the **Sample Mean** $\bar{X}_n$ as the arithmetic mean of these n observations:

$$\bar{X}_n = \frac{1}{n} \sum_{i=1}^n X_i$$

The Central Limit Theorem states that, as the sample size approaches infinity ($n \rightarrow \infty$), the probability distribution of $\bar{X}_n$ converges in distribution to a Normal distribution with the same mean as the population (μ) but with a reduced variance, equal to the original variance divided by the size of the sample (σ^2/n):

$$\bar{X}_n \xrightarrow{d} N\left(\mu, \frac{\sigma^2}{n}\right)$$

By standardizing the sample mean, the resulting Z statistic will converge to the Standard Normal distribution with a mean of 0 and a variance of 1:

$$Z = \frac{\bar{X}_n - \mu}{\sigma/\sqrt{n}} \xrightarrow{d} N(0, 1)$$

3 The Deep Meaning of the "Gaussian Collapse"

When it is said that *"every model, if enough examples are taken, collapses into a Gaussian model"*, it is fundamental to clarify a common misconception to avoid conceptual errors on the exam:

1. **The population does NOT change shape:** If the height of trees in a forest or factory defects do not follow a Gaussian distribution, they will remain non-Gaussian forever, even if we measure billions of trees. The raw data does not change its nature.
2. **What "collapses" is the MEAN of the samples:** If we take 100 trees at random and calculate the average height, then take another 100 trees and recalculate the average, and repeat this process thousands of times, it is **these average values** that will arrange themselves perfectly along a bell-shaped curve around the true mean μ .
(we will have few samples that will have a very low or very high mean given that it is highly improbable that the sample mean deviates significantly from the theoretical mean)

3.1 Why does the variance shrink? (σ^2/n)

From the formula, we notice that the variance of the line of means is $\frac{\sigma^2}{n}$. This means that as n increases (the size of the individual sample), the bell curve not only becomes more precise, but **it narrows and rises around μ** .

Visually, this happens because in a large sample (e.g., $n = 100$), the extremes compensate for each other: it is statistically almost impossible to pick 100 trees by pure chance that are all microscopically dwarf or all gigantic. The vast majority of 100-element samples will have a mean close to the central value. The larger the sample, the more the sample mean "locks" onto the actual real expected value of the population.

4 Graphical Visualization of Convergence

The graph below shows the behavior of the sample mean starting from a highly asymmetrical population (red curve, $n = 1$). As the sample size increases ($n = 5$ green curve, $n = 30$ blue curve), the distribution completely loses its original asymmetry and transforms into a perfect Gaussian bell curve centered on the real mean.




Figure 1: Visual demonstration of the Central Limit Theorem. Regardless of the original shape (red line), the distribution of the sample means collapses into a symmetrical bell curve (blue line) centered on μ .

5 Practical Example: Rolling Dice

Let us take the clearest example: rolling a 6-sided die. The population distribution is **Discrete Uniform**: every number from 1 to 6 has exactly the same probability of coming out ($1/6 \approx 16.6\%$). The probability graph is flat; it has nothing to do with a bell curve. The theoretical mean is $\mu = 3.5$.

- **If you roll $n = 1$ die at a time:** Your mean will simply be the number rolled. If you replicate the roll a thousand times, you will get an almost identical number of 1s, 2s, 3s, 4s, 5s, and 6s. The shape remains flat.
- **If you roll $n = 2$ dice at a time and take the average:** To get a mean of 1, you must roll a 1 and a 1 (only 1 combination). To get a mean of 3.5, you can roll (1,6), (2,5), (3,4), etc. (many more combinations). The distribution stops being flat and becomes triangular, with the peak on 3.5.
- **If you roll $n = 30$ dice at a time and take the average:** If you repeat this roll of thirty dice a thousand times and write down the averages obtained, you will discover that obtaining a mean of 1 or 6 is practically impossible. Almost all averages will concentrate around 3.5. If you draw the histogram of these means, you will see a perfect Gaussian curve appear.

6 What Is It Used For in Practice?

Without the Central Limit Theorem, modern inferential statistics simply would not exist. Its utility can be summarized in three crucial points:

1. **It allows testing and confidence intervals on unknown populations:** If you have to perform a quality control check in a multinational company, you do not need to know what mathematical law the production follows. You just need to extract sufficiently large samples ($n > 30$) and you know mathematically that the mean of those samples will obey the Gaussian law. You can therefore use Z -tests and Student's t -tests with absolute confidence.
2. **In the Real World:** Physical measurement errors, human heights, or IQs distribute themselves like a Gaussian. Each of these phenomena is the result of the **sum** of many small independent and random factors. And the sum of independent factors, by the CLT, collapses into a Gaussian.
3. **Computational Simplification:** It transforms potentially complex or intractable problems from a probability calculation perspective into simple calculations based on the Standard Normal distribution table.