

Detailed and Updated Summary:
*Towards leveraging the role of machine learning
and artificial intelligence in precision agriculture
and smart farming*
Computers and Electronics in Agriculture 198
(2022) 107119

NECKER

Based on the original PDF + some clarifications

December 7, 2025

Abstract

This document provides a technical and in-depth examination of the article "*Towards leveraging the role of machine learning and artificial intelligence in precision agriculture and smart farming*", published in *Computers and Electronics in Agriculture*. It explores the impact of digitalization in the agricultural sector, analyzing the Artificial Intelligence (AI) and Machine Learning (ML) methodologies applied to pre-harvest phases, crop management, irrigation, and yield prediction, as well as open challenges and future perspectives.

Contents

1	Introduction: The Agriculture 4.0 Revolution	3
2	Enabling Technologies: AI and Machine Learning	3
2.1	Supervised Machine Learning	3
2.2	Unsupervised Machine Learning	8
2.2.1	K-means:	8
2.2.2	Example: Agricultural Zonation via K-means	8
2.2.3	Fuzzy C-means (FCM):	9
2.2.4	Example: Soil Transition Mapping via FCM	10
2.2.5	Artificial Neural Networks (ANN - Unsupervised):	10
2.2.6	Example: Hyperspectral Spectral Analysis via KSOM	11
2.3	Dimensionality Reduction	12
2.3.1	PCA (Principal Component Analysis):	12
2.3.2	LDA (Linear Discriminant Analysis):	13
2.4	Deep Learning (DL)	14

2.4.1	CNN (Convolutional Neural Networks):	14
2.4.2	RNN (Recurrent Neural Networks) and LSTM (Long Short-Term Memory):	15
3	Applications in Agricultural Automation (Pre-Harvest)	17
3.1	Soil and Seed Analysis	17
3.2	Disease Detection and Classification	17
3.3	Crop Phenotyping	17
3.4	Fertilizer and Pesticide Management	18
3.5	Weed Detection	18
3.6	Smart Irrigation	18
4	Systems, Platforms and Yield Prediction	19
4.1	Crop Yield Prediction	19
4.2	Drones (UAV) and WSN	19
5	Open Challenges and Limitations	19
6	Future Perspectives	20

1 Introduction: The Agriculture 4.0 Revolution

Agriculture accounts for 6.4% of global GDP and constitutes the primary source of livelihood for millions of people. However, global food demand is projected to grow by 70% by 2050. This scenario imposes a paradigm shift towards sustainable practices, given that current food production consumes resources at unsustainable rates (70% of freshwater is destined for agriculture).

The article introduces the concepts of **Smart Farming (SF)** and **Precision Agriculture (PA)** as technological responses to these challenges.

- **Smart Farming (SF):** The application of IoT (Internet of Things) and AI technologies in the cyber-physical management of farms.
- **Precision Agriculture (PA):** The science aimed at improving crop yields and supporting management decisions through sensors and analysis tools, optimizing the use of fertilizers and irrigation.

The goal is to transform agriculture into a data-driven industry ("Agriculture 4.0"), capable of generating, transferring, and processing data to increase efficiency and reduce environmental impact.


2 Enabling Technologies: AI and Machine Learning

The document categorizes the computational techniques used in agriculture into four main macro-areas: supervised learning, unsupervised learning, dimensionality reduction, and deep learning.

2.1 Supervised Machine Learning

In this category, the algorithm learns a function that maps an input to an output based on example input-output pairs provided during the training phase (for instance, thousands of images with a label indicating what the model should predict from that image).

- **Support Vector Machine (SVM):**




This classifier constructs a hyperplane (or a set of hyperplanes) in a high-dimensional space, aiming to maximize the "margin", which is the distance between the closest data points of the different classes.

- **Mechanism:** It uses the "kernel trick" to transform a low-dimensional input space into a higher-dimensional one, making linearly non-separable problems separable. Optimization is often handled via the SMO (Sequential Minimal Optimization) algorithm.
- **Advantages and Limitations:** It is particularly effective in high-dimensional spaces, even when the number of dimensions exceeds the number of samples. However, it is not scalable to very large datasets, and performance degrades in the presence of high noise or overlapping classes.

Example: Apple Classification via SVM

- * **Scenario:** Separating apples into **Premium** and **Standard** classes using measurements of *Diameter* and *Firmness*.
- * **Operation:** The SVM is not satisfied with just any separation line, but finds the one that guarantees the maximum tolerance margin, defined by the *Support Vectors* (the apples located right on the boundary between classes). If linear separation fails (e.g., Premium apples are nested in the center), the algorithm applies the **Kernel Trick** (e.g., RBF Kernel). This projects the features into a higher-dimensional space where the classes become linearly separable.

• Decision Trees:




A non-parametric model that recursively splits data into nodes to maximize discrimination between classes.

- **Split Criteria:** The split can be based on a single feature (*Univariate*) or on multiple features (*Multivariate*). To quantify the quality of the split, the Gini Index, Information Gain, and Entropy are used.
- **Characteristics:** It does not require data normalization or scaling and handles both numerical and categorical features well without hot-encoding.
- **Critical Issues:** It suffers from instability (small variations in data drastically alter the tree structure) and tends to overfit. Furthermore, it is inadequate for predicting continuous values in complex regression tasks.

Example: Apple Classification via Decision Trees

- * **Scenario:** Rapidly classifying apples as **Premium** or **Reject** for the elimination of defective fruits.
- * **Operation:** The algorithm creates a series of sequential questions. At the first node, it chooses the feature that reduces impurity the most (increases homogeneity): for example, "Is firmness ≤ 6 ". On the resulting branch, the next most informative question is asked, for example: "Is diameter ≤ 5 cm?".
- * **Output:** The result is a clear and interpretable decision path: **If Firmness ≤ 6 AND Diameter ≤ 5 , then CLASS = Reject.**

- **k-Nearest Neighbors (k-NN):**




Defined as a "lazy learning" or instance-based learning (Instance-Based Learning - IBK) method, k-NN does not construct a model during training but directly uses training instances to classify new data.

- **Operation:** It creates a locally optimized model for each test instance based on proximity to the k nearest neighbors.
- **Advantages and Limitations:** It is simple to implement for multi-class problems and allows adding new data without retraining. However, it is sensitive to outliers, requires homogeneous features, and suffers from the "curse of dimensionality", becoming slow with large datasets.

Example: Apple Classification via k-NN

- * **Scenario:** Classifying an apple with features ($x_1 = 7.2, x_2 = 6.5$) as 'Premium' or 'Standard'.
- * **Operation:** The algorithm does not have a model trained in advance. At the time of classification, it calculates the Euclidean distance of the new apple from **all** previously labeled apples. The k nearest neighbors are selected (e.g., $k = 5$). If 4 out of the 5 neighbors are labeled as 'Premium', the new apple is classified by majority vote as **Premium**.

- **Random Forest:**




An ensemble algorithm that aggregates multiple decision tree classifiers to improve overall accuracy and mitigate the shortcomings of a single tree.

- **Mechanism:** Each tree in the ensemble determines a class for the sample, and the final label is decided via a majority vote.
- **Performance:** It is robust to outliers and noise, reduces the risk of overfitting, and handles non-linear data well. On the other hand, training is slow, it can introduce bias with categorical variables, and it is not ideal for linear methods with many sparse features.

Example: Apple Classification via Random Forest

- * **Scenario:** High-precision classification of apple quality using multiple features (Diameter, Firmness, Weight, Acidity).
- * **Operation:** A "forest" of Decision Trees is created. Each tree is trained on a random bootstrap sample of the data and, crucially, on a random subset of features (ensuring decorrelation between trees). For the final classification of a new apple, each tree casts a vote ('Premium' or 'Standard'). The majority vote of the forest is the final prediction.

- **Naïve Bayes:**



A probabilistic generative model based on Bayes' theorem, known for its speed and ability to handle high-dimensional data.

- **Assumptions:** It relies on the (often unrealistic) assumption that all features are independent of each other.
- **"Zero Frequency" Problem:** If a categorical variable not present in the training set appears in the test dataset, the model assigns it a zero probability, making prediction impossible (a phenomenon known as "Zero Frequency"). It is nonetheless excellent for categorical values and multi-class prediction problems.

Example: Apple Disease Detection via Naïve Bayes

- * **Scenario:** Predicting the probability that an apple is ****Diseased**** given binary features (e.g., Presence of Spots X_1 , Abnormal Color X_2).
- * **Operation:** The algorithm calculates the posterior probability of the apple being diseased (Class $Y = 1$) by multiplying the prior probability of Y by the conditional probabilities of each individual feature, assuming that $P(X_1|Y) \times P(X_2|Y)$ holds true.
- * **Conceptual Formula:** $P(\text{Diseased}|X_1, X_2) \propto P(\text{Diseased}) \times P(X_1|\text{Diseased}) \times P(X_2|\text{Diseased})$.

2.2 Unsupervised Machine Learning

Unsupervised Machine Learning applies when the input data lacks predefined class labels. Its primary objective is to explore the intrinsic structure of the dataset, identifying **hidden patterns**, **natural groupings** (clustering), or **latent structures** (dimensionality reduction) for diagnostic and exploratory purposes in precision agriculture (in this case, we no longer have labels telling us, for example, whether an apple is premium or not).

2.2.1 K-means:


K-means operates in an iterative loop that aims to minimize the **Sum of Squared Errors (SSE)**, which is the total squared distance between points and the centroids of their respective clusters. The process consists of three main phases:

1. **Initialization:** k points are randomly chosen as initial centroids (μ_i).
2. **Assignment (E-Step):** Each data point is assigned to the cluster whose centroid is **geometrically closest** (typically measured with Euclidean distance).
3. **Update (M-Step):** Centroids are recalculated as the **arithmetic mean** of all points assigned to the cluster.

The cycle repeats (steps 2 and 3) until the centroids stabilize or convergence is reached.

Limitations

It is efficient but heavily sensitive to the initial choice of k and centroids. Furthermore, the presence of **outliers** can significantly distort the centroids, leading to sub-optimal groupings.



2.2.2 Example: Agricultural Zonation via K-means

- **Scenario:** Subdividing a large agricultural field into homogeneous regions (management zones) based on soil properties or vegetative vigor.
- **Mechanism:** The algorithm requires the user to specify the number of clusters k in advance (for example, 3 zones: High, Medium, Low Vigor). Initially, k random centroids are chosen. Each data point (e.g., soil measurement) is assigned to the nearest centroid based on Euclidean distance. The centroids are recalculated as the mean of all assigned points, and the process repeats until the centroids converge (no longer change position).

2.2.3 Fuzzy C-means (FCM):

The objective of FCM is to minimize a cost function (called the objective function) which is a variation of the SSE, weighted by membership degrees u_{ij} . The process is iterative:

1. **Initialization:** The number of clusters (C) and the fuzziness exponent m (typically $m = 2$) are established. The **Membership Matrix** U is randomly initialized, where element u_{ij} represents the membership degree of point i to cluster j . The sum of the membership degrees for any point must equal 1 ($\sum_{j=1}^C u_{ij} = 1$).
2. **Center Update:** Centroids μ_j are calculated as a weighted average of the data points. The weight of each point is given by its membership degree raised to the fuzziness exponent m :

$$\mu_j = \frac{\sum_{i=1}^N (u_{ij})^m \cdot x_i}{\sum_{i=1}^N (u_{ij})^m}$$


3. **Membership Matrix (U) Update:** The matrix U is updated based on the distance of each point from the new centroids, strengthening membership in closer clusters and diminishing it for farther ones:

$$u_{ij} = \frac{1}{\sum_{k=1}^C \left(\frac{\|x_i - \mu_j\|}{\|x_i - \mu_k\|} \right)^{\frac{2}{m-1}}}$$

The cycle repeats until convergence is reached or when the maximum variation in membership degrees falls below a minimum threshold.

Key Advantage

The "fuzzy" nature of FCM is fundamental in precision agriculture, especially when analyzing data with **ambiguous boundaries**, such as the transition between two soil types or the gradual variation of water stress in a field. FCM handles overlapping data better and offers a more granular and robust view of transition zones compared to hard clustering.



2.2.4 Example: Soil Transition Mapping via FCM

- **Scenario:** Analyzing soil moisture data to identify regions where the transition between "dry" and "wet" is not sharp.
- **Mechanism:** Unlike K-means, a data point does not belong rigidly to only one cluster, but has a **membership degree** to all C clusters, quantified by a membership matrix. If an area has intermediate moisture, the data point might have a 60% membership in the 'Wet' cluster and 40% in the 'Intermediate' cluster.

2.2.5 Artificial Neural Networks (ANN - Unsupervised):

Unsupervised ANNs are a class of models, distinct from supervised networks (like CNNs), that do not require training labels. They are primarily used for feature learning; among these, KSOMs (or Self-Organizing Maps, SOM) are a type of neural network that implements unsupervised competitive learning to produce a low-dimensional representation (usually 2D) of input vectors, while maintaining the topological and proximity relationships of the original data.

1. **Structure:** The KSOM consists of an **input layer** and a **map layer** (typically a 2D grid of neurons). Each neuron on the map is connected to all neurons in the input layer via a weight vector.
2. **Competitive Mechanism:** Given an input vector x (e.g., the spectral profile of a pixel), the algorithm calculates the distance between x and the weight vector w_j of each neuron j on the map. The neuron with the closest weight to the input is declared the **Best Matching Unit (BMU)**.


$$BMU = \arg \min_j \|x - w_j\|$$

3. **Topological Update:** Unlike K-means, not only the BMU is updated, but also all neurons within its **topological neighborhood** on the 2D grid. The amount

of update is inversely proportional to the distance from the BMU on the map (the further away, the less it is updated) and decreases over time (the neighborhood radius gradually shrinks). This ensures that similar data are mapped onto neighboring neurons, preserving structure.

Key Advantage in Agriculture

The KSOM is crucial for analyzing extremely complex, high-dimensional agricultural data, such as **hyperspectral** images (with hundreds of bands). It allows visualizing and clustering similar spectral profiles (which could indicate the same water stress level or disease) on a simple 2D map, enabling agronomists to identify latent patterns and stress groupings without having to label the data beforehand.



2.2.6 Example: Hyperspectral Spectral Analysis via KSOM

- **Scenario:** Analyzing very high-dimensional hyperspectral images (hundreds of spectral bands per pixel) to identify vegetation types or water stress without initial labels.
- **Mechanism (KSOM):** The KSOM is a type of neural network that creates a low-dimensional map (often 2D) that preserves the topological relationships of high-dimensional input data. Neighboring nodes on the 2D map represent similar input vectors. In an agricultural context, if two pixels have similar spectral profiles (e.g., both represent a healthy crop), they will be mapped to neighboring cells in the KSOM map, creating visual clusters.

2.3 Dimensionality Reduction


Dimensionality Reduction is an essential pre-processing phase in Machine Learning, particularly for agricultural data (such as multispectral or hyperspectral measurements) characterized by a high number of variables ($\mathbf{X}$). The goal is to transform the dataset from a high-dimensional space to a low-dimensional one, **preserving as much information as possible** or maximizing the **separation between classes**. This process reduces computational complexity and mitigates the effect of the "curse of dimensionality" (a phenomenon where data becomes too sparse in very large spaces).

2.3.1 PCA (Principal Component Analysis):

PCA is an **unsupervised dimensionality reduction technique** that finds linear subspaces that maximize the variance of the projected data.

Mechanism and Application (PCA)

- **Objective:** To identify a new set of orthogonal variables, called **Principal Components** (PCs), which are linear combinations of the original variables. The first PC captures the maximum possible variance in the data, the second PC captures the maximum residual variance, and so on.
- **Operation:** The algorithm calculates the eigenvectors and eigenvalues of the data's covariance matrix. The PCs are the eigenvectors, ordered by the magnitude of their eigenvalues (which represent the explained variance). Only the first k PCs that explain a significant percentage of the total variance (e.g., 95%) are selected.
- **Agricultural Benefit:** It is used for the **compression** of multispectral or hyperspectral images. For example, from 100 spectral bands (features), 5 or 10 Principal Components can be selected that contain almost all relevant information about plant health, drastically reducing training time for subsequent classifiers (such as SVM or CNN).



2.3.2 LDA (Linear Discriminant Analysis):


LDA is a **supervised dimensionality reduction technique** specifically designed for classification problems.

Mechanism and Application (LDA)

- **Objective:** To project data into a lower-dimensional space that **maximizes the separation between classes** (inter-class variance) and, simultaneously, **minimizes variance within each class** (intra-class variance).
- **Operation:** Unlike PCA (which ignores labels), LDA uses the Y labels to find the best projective direction. The algorithm calculates a measure of separability and searches for projections that maximize the ratio between inter-class variance and intra-class variance.
- **Agricultural Benefit:** If the goal is to distinguish between two classes (e.g., "Healthy Plant" vs. "Diseased Plant") using data with many features (e.g., 20 vegetation indices), LDA ensures that the projected features are the most effective for classification. LDA is often used as a pre-processor before applying simple classifiers.

Fundamental Difference: Objective

While **PCA** is an unsupervised method that aims to **preserve as much information as possible** (maximizing total variance), **LDA** is a supervised method that aims to **maximize class discrimination** (optimizing the separation between classes). For a classification problem, if labels are available and reliable, LDA is generally more powerful.



2.4 Deep Learning (DL)


Deep Learning (DL) represents the evolutionary pinnacle of Artificial Neural Networks (ANNs) and is distinguished by using architectures with **multiple hidden layers** ("deep"). Its key feature is the elimination of the need for **manual feature engineering**: the model is able to automatically learn and extract the hierarchy of relevant features directly from raw data (unlike **unsupervised ML**, here we have labels but the model must extract the features).

2.4.1 CNN (Convolutional Neural Networks):

CNNs are feed-forward neural network architectures highly specialized for the analysis and processing of grid-structured data, such as images and spectral data. They are the dominant technology for computer vision in agriculture.

Mechanism and Application (CNN)

- **Convolutional Layer:** The central function is the **convolutional filter** (kernel) that slides over the input image (or previous activation maps), performing element-by-element multiplication to extract local features (edges, corners, textures). Each subsequent layer learns features of increasing complexity.
- **Pooling Layer:** Reduces spatial dimensionality (compression), decreasing the number of parameters and making the extracted features less sensitive to small translations in the input.
- **ReLU Activation Function:** (Rectified Linear Unit) Introduces non-linearity, allowing the network to learn complex relationships that would otherwise be impossible with linear transformations.
- **Agricultural Application:** They are employed for **disease recognition** (e.g., leaf classification in the PlantVillage dataset), **fruit detection** (for automated harvesting), and identification of **weeds** (weed detection) via drones. Pre-trained architectures like AlexNet, VGG, and ResNet are often reused via **Transfer Learning**.




2.4.2 RNN (Recurrent Neural Networks) and LSTM (Long Short-Term Memory):

RNNs are designed specifically to process **sequential data** or **time series** (time, pixel sequence, words). Unlike feed-forward networks, they maintain a "memory" of past inputs.

Mechanism and Application (RNN/LSTM)

- **The Limitation of Classical RNNs:** Standard RNNs suffer from the **vanishing gradient** problem, which prevents them from learning long-term dependencies in the data (e.g., the correlation between rainfall 6 months ago and current yield).
- **LSTM (Long Short-Term Memory):** This is the most widely used variant and solves the problem by introducing **memory cells** and **gates** (Input, Forget, Output). These gates selectively control what information should be kept in long-term memory, what should be forgotten, and what should be exposed.
- **Agricultural Application:** LSTMs are crucial for crop **yield prediction**. They can effectively model the temporal dependence between environmental variables (historical climate data, temperatures, precipitation, vegetation indices measured over time) and final crop yield.



Summary Table:

Methodology	Mode	Feature Source (X)	Label Required (Y)
Traditional ML	Supervised (e.g., SVM)	Human-provided (Manual Feature Engineering)	YES
	Unsupervised (e.g., K-means)	Human-provided (Manual Feature Engineering)	NO
Deep Learning	Supervised (e.g., CNN)	Automatically Learned (from raw data)	YES
	Unsupervised (e.g., Autoencoder)	Automatically Learned (from raw data)	NO

3 Applications in Agricultural Automation (Pre-Harvest)

The article analyzes in detail how these technologies are applied in the phases prior to harvesting.

3.1 Soil and Seed Analysis

Classification of soil properties helps reduce fertilizer costs. Models such as ELM (Extreme Learning Machine) and SVM are used to predict pH, organic carbon, and moisture. For seeds, computer vision and CNNs (e.g., VGG19, GoogLeNet) are used to:

- Classify high-quality vs. low-quality seeds (e.g., in pepper).
- Estimate the number of seeds per pod (e.g., soybean).
- Identify morphological defects.

3.2 Disease Detection and Classification

This is one of the most critical areas. Early identification prevents massive losses.

- **Conventional Methods:** They utilize manually extracted features (color, texture) and classifiers like SVM or k-NN. Examples include the detection of fungal diseases on grapes, citrus, and soybeans with accuracies ranging from 88% to 97%.
- **Deep Learning:** CNNs (AlexNet, ResNet, VGG) have shown superior performance, reaching accuracies over 99% (e.g., PlantVillage dataset). DL handles variability better and does not require manual feature engineering, making it suitable for diagnostic mobile apps (e.g., Plantix).

3.3 Crop Phenotyping

Automated phenotyping measures physical and biochemical characteristics of plants (growth, stress tolerance). Systems based on hyperspectral sensors, LiDAR, and autonomous vehicles (ground robots or drones) allow overcoming the "bottleneck" of manual measurements, accelerating genetic selection programs.

3.4 Fertilizer and Pesticide Management

Excessive use of agrochemicals damages the ecosystem and human health.

- **Pest Detection:** Computer vision and acoustic models identify specific harmful insects, allowing targeted application of pesticides (e.g., "See & Spray" technology that reduces chemical use by 80%).
- **Optimization:** Predictive algorithms suggest the exact amounts of nutrients needed based on real-time soil conditions.

3.5 Weed Detection

Weeds compete for resources. The article highlights the use of drones and ground robots equipped with multispectral cameras and algorithms (e.g., Random Forest, CNN) to distinguish crops from weeds, allowing localized treatments (spot spraying) instead of broadcast application.

3.6 Smart Irrigation

Water optimization is vital. IoT systems with soil moisture sensors and weather stations, combined with algorithms like Naïve Bayes or decision trees, automatically control irrigation valves. Commercial platforms (e.g., ConsertWater) use satellite data to predict water requirements, saving up to 30% of water.

4 Systems, Platforms and Yield Prediction

4.1 Crop Yield Prediction

Predicting how much a field will produce is essential for food security and the economy.

- It is a complex problem influenced by genotype, environment, and management.
- Deep Learning (particularly LSTM and hybrid CNN-RNN architectures) outperforms traditional methods because it successfully models the non-linear relationships between historical data, weather, and satellite imagery.
- Cited examples include predicting soybean yield in Brazil and rice yield in Korea, integrating spatial and temporal data.

4.2 Drones (UAV) and WSN

Drones are not just monitoring tools but active platforms for mapping plant vigor (NDVI) and spraying. Wireless sensor networks (WSN) form the core infrastructure for continuous data collection (temperature, humidity, pH), transmitting via protocols like LoRaWAN or ZigBee.

5 Open Challenges and Limitations

Despite the progress, the document identifies several significant barriers:

1. **Data Scarcity and Quality:** Lack of large public labeled datasets (beyond PlantVillage). Many datasets are small or proprietary.
2. **Real-World vs. Laboratory Conditions:** Most models are trained on images with uniform backgrounds and controlled lighting, failing in the open field where there are shadows, overlaps, and complex backgrounds.
3. **Standardization:** Interoperability between different devices and data formats is lacking.
4. **Scalability:** Handling heterogeneous "Big Data" requires scalable algorithms and powerful hardware (GPU/TPU).
5. **Drones:** Battery limitations (autonomy < 1 hour), high costs of hyperspectral sensors, and restrictive flight regulations.

6 Future Perspectives

The article concludes by suggesting priority research directions:

- **Identification of Disease Stage:** It is not enough to detect the disease; it is necessary to understand what stage it is in to decide on treatment.
- **Damage Quantification:** Developing systems that calculate the percentage of infection or damage to dose chemicals.
- **Offline Mobile Applications:** Developing lightweight and efficient apps that farmers can use directly in the field, even without an internet connection.
- **Transfer Learning:** Utilizing pre-trained models to mitigate the shortage of specific data for certain crops.

References

- [1] Tawseef Ayoub Shaikh, Tabasum Rasool, Faisal Rasheed Lone (2022).
Towards leveraging the role of machine learning and artificial intelligence in precision agriculture and smart farming.
Computers and Electronics in Agriculture, 198, 107119.
<https://doi.org/10.1016/j.compag.2022.107119>